

1

```
cities <- data.frame(ran <-runif (30,min=25,max=40))
```

```
i <- 1
```

```
while(i<10){
```

```
  ran <-runif (30,min=25,max=40)
```

```
  cities = cbind(cities, ran)
```

```
  i <- i+1
```

```
}
```

```
names(cities) <- c("c1","c2","c3","c4","c5","c6","c7","c8","c9","c10")
```

```
#average of every column, manual calculation
```

```
sum <- function (){
```

```
  sum <- colSums(cities)/30
```

```
  return(sum)
```

```
}
```

```
sum()
```

#manual function to get minimum number on every column.

```
minvalue <- sort(cities$c1)
```

```
min1 <- minvalue[1]
```

```
minvalue <- sort(cities$c2)
```

```
min2 <- minvalue[1]
```

```
minvalue <- sort(cities$c3)
```

```
min3 <- minvalue[1]
```

```
minvalue <- sort(cities$c4)
```

```
min4 <- minvalue[1]
```

```
minvalue <- sort(cities$c5)
```

```
min5 <- minvalue[1]
```

```
minvalue <- sort(cities$c6)
```

```
min6 <- minvalue[1]
```

```
minvalue <- sort(cities$c7)
```

```
min7 <- minvalue[1]
```

```
minvalue <- sort(cities$c8)
```

```
min8 <- minvalue[1]
```

```
minvalue <- sort(cities$c9)
```

```
min9 <- minvalue[1]
```

```
minvalue <- sort(cities$c10)
```

```
min10 <- minvalue[1]
```

```
#sorting every column so the value of of the lowest temperature
```

```
#is placing itself on position 1.
```

```
mintot <- data.frame(min1,min2,min3,min4,min5,min6,min7,min8,min9,min10)
```

```
#Printing the vectors of every position 1
```

```
View(mintot)
```

```
#manual function to get maximum number on every column.
```

```
maxvalue <- sort(cities$c1,decreasing = TRUE)
```

```
max1 <- maxvalue[1]
```

```
maxvalue <- sort(cities$c2,decreasing = TRUE)
```

```
max2 <- maxvalue[1]
```

```
maxvalue <- sort(cities$c3,decreasing = TRUE)
```

```
max3 <- maxvalue[1]
```

```
maxvalue <- sort(cities$c4,decreasing = TRUE)
```

```
max4 <- maxvalue[1]
```

```
maxvalue <- sort(cities$c5,decreasing = TRUE)
```

```
max5 <- maxvalue[1]
```

```
maxvalue <- sort(cities$c6,decreasing = TRUE)
```

```
max6 <- maxvalue[1]
```

```
maxvalue <- sort(cities$c7,decreasing = TRUE)
```

```
max7 <- maxvalue[1]
```

```
maxvalue <- sort(cities$c8,decreasing = TRUE)
```

```
max8 <- maxvalue[1]
```

```
maxvalue <- sort(cities$c9,decreasing = TRUE)
```

```
max9 <- maxvalue[1]
```

```
maxvalue <- sort(cities$c10,decreasing = TRUE)
```

```
max10 <- maxvalue[1]
```

```
#sorting every column backwards so the value of of the highest temperature
```

```
#is placing itself on position 1.
```

```
maxtot <- data.frame(max1,max2,max3,max4,max5,max6,max7,max8,max9,max10)
```

```
#Printing the vectors of every position 1
```

```
View(maxtot)
```

```
#inbuild function to get average of every column
```

```
for (i in colnames(cities)){
```

```
  print(mean(cities[[i]]))
```

```
}
```

```
#inbuild function to get minimum number on every column
```

```
minimum <- min(cities)
```

```
return(minimum)
```

```
#inbuild function to get maximum number on every column
```

```
maximum <- max(cities)
```

```
return(maximum)
```

2

```
library(tidyverse)

kidney <- read.csv("C:/Data Scientist/R programming/Final Assignment/kidney_disease.csv")

class(kidney)

#Checking the structure of data and coming up with the list of problems.

View(kidney)

str(kidney) #to see the setup and what needs to be changed. All the class
#for every column. These will be referred to in every column hereafter.

(sapply(kidney, unique))#gives unique values, to identify what is more than just
#an empty value, NA, or/and misspelled.

#Column "id" is a copy of the ruler of the y axis, with a slight deviation,
#since it starts at "0". It will be removed.

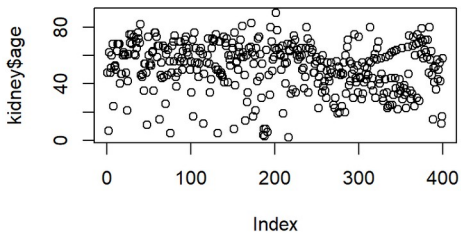
kidney$id <- NULL #This removed the column "id".

#Column "age"

#to see that there is 9 NA in this column.

#Plotting the "age" column shows that age on the participants is widespread and
#an average would be suited to #use in this case.

plot(kidney$age)
```



```
kidney$age[is.na(kidney$age)] <- (mean(kidney$age,
na.rm = TRUE)) # to change

#all the NA to the #average(mean) number. The average
is 51.4725, which means it

#has decimals that the rest of the column doesn't have.
#The column is also set

#to numeric. Since this is an average, it will comfort to
the rest of the column


#and be forced to be a number without decimal.

kidney$age <- as.integer(kidney$age)

#Column "bp" have all its positions in the tenth of number. That is 20,30,40 and
#so on. It also have 12 NA out of 400 #rows. The plot shows a widespread of numbers.
#So a "mean" could replace the positions.

mean(kidney$bp[!is.na(kidney$bp)]) # the mean #calculated to 76.46907. This will
```

#be rounded up to 80 #for 6 of the positions and 70 for the other 6, so not to
#mess up then actual settings of whole tenth numbers.



kidney[8, 2] = 80

kidney[76, 2] = 80

kidney[133, 2] = 80

kidney[139, 2] = 80

kidney[162, 2] = 80

kidney[165, 2] = 80

#The rest is NA, so they will be changed to 70 with

kidney\$bp[is.na(kidney\$bp)] <- 70

#in "sg", "al" and "su", the NA, is almost all at the same rows. In "sg",

#47 values are set to NA. Reading the Kaggle info, "sg" value is Specific Gravity

#and is the relative density between one substance and a reference substance.

#The range of this value is between 1,005 to 1,025. Without any reference point

#of the magnitude of these values they will not get an average value due to their

#slim variation and comparing to the index in the picture below, any deviation

#might give a less accurate plot. A KNN imputation would give a better estimated

#value. But #since a KNN is not available, these rows will be deleted so they won't

#obscure the rest of the table.

which(is.na(kidney\$sg)) # show all the NA in column "sg" to visually see how it

#correspond to "al" and "su".

#They have their own NA that "sg" doesn't. This will be dealt with separately.

kidney <- kidney [-c(14,18,22,29,31,38,51,58,60,79,82,83,87,99,100,105,107,110,

122,123,126,133,143,149,152,161,163,167,175,189,198,202,203,

204,206,209,223,229,232,237,239,246,269,281,296,323,347),]

#These rows are for "sg" and with the code above they were deleted.

rownames(kidney) <- 1:353# to have the accurate row numbers.

#"al" and "su" column is reset to integer

kidney\$al <- as.integer(kidney\$al)

kidney\$su <- as.integer(kidney\$su)

#The rows that "al" and "su" have NA separately is

```

kidney$al[is.na(kidney$al)]

#Svar 2

#and

kidney$su[is.na(kidney$su)]

#Svar 4

#They will also be removed

which(is.na(kidney$al))

#Svar 113, 115

which(is.na(kidney$su))

#Svar 74, 113, 115, 165

kidney <- kidney [-c(74, 113, 115, 165),]

rownames(kidney) <- 1:349 #To reset the rownames.

#Column "rbc" is a chr, and it seem to have a lot of empty spaces. To see how many,

#kidneytest is created

kidneytest <- kidney

kidneytest$rbc [kidneytest$rbc == ''] <- NA# and all the empty spaces become NA.

kidneytest$rbc[is.na(kidneytest$rbc)] # then count the NA in kidneytest that

#actually accounts for the empty #values in kidney. 107 of the cells are empty.

#This column together with other columns farther down the row-lane has equally

#missing values and NA. This column will be untouched in this phase.

#With the same method as in "rbc", column"pc" will make use of kidneytest since

#it has no NA and only empty values:

kidneytest$pc [kidneytest$pc == ''] <- NA

kidneytest$pc[is.na(kidneytest$pc)]

# the amount of NA in this column is 21. The issue here is that being chr,

#and the values are binary, a not answer could have a certain value. The empty

#spaces will be set to NA. The column will be changed to a factor with value 1

#for "normal", 2 for "abnormal" and 3 for "NA".

kidney$pc [kidney$pc == ''] <- NA

#The column "pc" is set to chr, so it has to be converted to a factor.

kidney$pc <- as.factor(kidney$pc)

```

```

kidney$pc <- factor(kidney$pc, levels = c("normal", "abnormal", "NA"))
#The column "pcc" is also a chr, and has the unique values of "present",
#"notpresent" and empty spaces.
#To check how many spaces, kidneytest is used again.
kidneytest <- kidney
kidneytest$pcc[kidneytest$pcc == ""] <- NA
kidneytest$pcc[is.na(kidneytest$pcc)] # To count how many empty spaces kidney$pcc
#have, and its 4 of them.
#This column will be treated as the previous one, where the empty spaces will be
#filled with NA, and the factor will be set to 1 as "present", 2 as "notpresent"
#and 3 as "NA". These are the steps
kidney$pcc[kidney$pcc == ""] <- NA
kidney$pcc <- as.factor(kidney$pcc) # to make the column a factor.
kidney$pcc <- factor(kidney$pcc, levels = c("present", "notpresent", "NA"))
#The column "ba" is also a chr. I have the same setup as "pcc", with only 4 empty
#spaces. So, the same application will be applied to "ba" as in "pcc". The steps
#are summarized below.
kidneytest$ba[kidneytest$ba==""] <- NA
kidneytest$ba[is.na(kidneytest$ba)] #This was the test to determine the amount of NA.
unique(kidney$ba) # to determinate the different value names. Those are "present",
#"notpresent" and empty value.
kidney$ba[kidney$ba == ""] <- NA
kidney$ba <- as.factor(kidney$ba)
kidney$ba <- factor(kidney$ba, levels = c("present", "notpresent", "NA"))
#Column "bgr" is numeric with NA values.
which(is.na(kidney$bgr)) # to show that bgr have 38 NA values. This is more than
#1/10 of all the values since rows has been deleted and is down to 349. Using the
#kidneytest dataframe to see how many NA were in bgr from the start gives
which(is.na(kidneytest$bgr)) # 40, so it is still 1/10 of the rows without a value.
#The initial rows were 400.
#It would be difficult to create an average to mimic the rest of the values.

```



```

#This is sort of trend for the 5 columns after, "bu", "c", "sod", "pot" and "hemo".
#They are all numeric, and they all have a large chunk of NA.
which(is.na(kidney$bu)) # with 18 NA.
which(is.na(kidney$sc)) # with 17 NA.
which(is.na(kidney$sod)) # with 75 NA.
which(is.na(kidney$pot)) # with 75 NA.
which(is.na(kidney$hemo)) # with 39 NA.
# The NA together with the kidneytest of "rbc", that have 107 NA, there are similar
#rows. Doing a visual check, it seems to correspond well. Having these nonvalues
#in the dataset will only obscure the registered values.
#The rows of "rbc" in test version will be deleted first to make this estimation more solid.
kidneytest$rbc[kidneytest$rbc==""] <- NA # to rbc empty values to NA
which(is.na(kidneytest$rbc)) # These are the rows that will be deleted.
kidneytest <- kidneytest [-c(1,2,6,7,11,13,15,16,17,18,21,23,24,27,28,36,37,40,
41,42,45,46,47,52,54,56,57,60,61,62,64,67,69,70,71,
75,76,81,82,83,84,86,87,89,91,92,93,95,97,99,100,
101,102,103,104,105,108,110,112,114,115,116,125,129,
134,135,136,138,139,140,142,144,148,151,152,153,154,
160,165,166,168,169,170,172,175,177,178,179,180,181,
183,184,187,193,194,196,197,199,200,243,261,299,300,
331),]
rownames(kidneytest) <- 1:245 # This is done in kidneytest to see if the the actual
#similarities to the other colums
which(is.na(kidneytest$bu)) # with 12 NA.
which(is.na(kidneytest$sc)) # with 11 NA.
which(is.na(kidneytest$sod)) # with 34 NA.
which(is.na(kidneytest$pot)) # with 34 NA.
which(is.na(kidneytest$hemo)) #with 22 NA.
# So the corresponding NA to the column "bgr" is almost half compared. Had this
#been a lower number this wouldn't have been significant, but somehow it has a
#connection. So the test will be made real on kidney dataset, and the rows for

```

#“bu”, “sc”, “sod”, “pot” and “hemo” will be deleted as well.

```
kidney$rbc[kidney$rbc==""] <- NA
```

```
kidney <- kidney [-c(1,2,6,7,11,13,15,16,17,18,21,23,24,27,28,36,37,40,41,42,45,  
46,47,52,54,56,57,60,61,62,64,67,69,70,71,75,76,81,82,83,84,  
86,87,89,91,92,93,95,97,99,100,101,102,103,104,105,108,110,  
112,114,115,116,125,129,134,135,136,138,139,140,142,144,148,  
151,152,153,154,160,165,166,168,169,170,172,175,177,178,179,  
180,181,183,184,187,193,194,196,197,199,200,243,261,299,300,  
331),]
```

rownames(kidney) <- 1:245 # kidney is now the same as the test.

#The remaining in column “rbc” will be changed to the class factor and the

unique(kidney\$rbc) # shows the only 2 values left are “normal” and abnormal

#The factors will be set as 1 as “normal and 2 as ”abnormal”.

```
kidney$rbc <- as.factor(kidney$rbc)
```

```
kidney$rbc <- factor(kidney$rbc,levels=c("normal","abnormal"))
```

#Column “bu”

```
which(is.na(kidney$bu)) # to have the rows to get deleted.
```

```
kidney <- kidney[-c(25,26,32,69,91,97,131,137,163,166,184,225),]
```

```
rownames(kidney) <- 1:233
```

#Column “sc”

```
which(is.na(kidney$sc)) # This is now 0 NA.
```

#Column “sod”

```
which(is.na(kidney$sod))
```

```
kidney <- kidney[-c(1,3,4,5,13,14,15,16,20,24,36,42,44,45,46,58,59,62,66,70,72,  
75,93,147,175,199),]
```

```
rownames(kidney) <- 1:207 # The numbers of NA, is shrinking steady, so the NA for
```

#the 5 column seem to connect on some level.

#Column “pot” has no 0 NA

```
which(is.na(kidney$pot))
```

#Column “hemo”

```
which(is.na(kidney$hemo))
```

```

kidney <- kidney[-c(17,19,24,30,38,41,53,65,69,70,74,99,136,140,144,146,175),]
rownames(kidney) <- 1:190

#Column "pcv" has no NA, but there some values missing. It is also a chr, but
#with numbers.

(unique(kidney$pcv)) # shows that there are one position that is missing a value.
#This position will be replaced by an average.

kidney$pcv[kidney$pcv==""] <- NA # The blank space is first changed to NA
kidney$pcv <- as.integer(kidney$pcv) # the class is changed to integer so we can
#perform a mean function.

kidney$pcv[is.na(kidney$pcv)] <- (mean(kidney$pcv, na.rm = TRUE)) # The NA is now
#the mean. The problem is that the mean is a decimal value. So, even though it is
#an integer, it shows decimal. The new value is forced into no decimal value with
#another

kidney$pcv <- as.integer(kidney$pcv)

#Column "wc" is shown to have errors in spacing, that is \t before a value, and
#empty values shown from

unique(kidney$wc)

kidney$wc [is.na(kidney$wc)] # it has 0 NA


#Doing a kidneytest to convert the empty spaces in "wc" to NA to see how many give us

kidneytest <- kidney

kidneytest$wc[kidneytest$wc == ""] <- NA

kidneytest$wc [is.na(kidneytest$wc)] # There are 12 NA.

```



#Plotting this test to see the variation, one can determinate that it has a spread

#and a mean to the blank spaces would be preferable.

#But first this needs to be changed to an integer since it is set to chr.

```
kidney$wc <- as.integer(kidney$wc)
```

#mean(kidney\$wc, na.rm = TRUE)) # The average value of wc is 8457.865. This will

#be rounded not to make unique values when plotting, to suit the style of other

#values in "wc" = 8500

```
kidney$wc[is.na(kidney$wc)] <- 8500
```

#Column "rc"

#Using kidneytest to determinate how many NA the column of "rc" have gives:

```
kidneytest <- kidney
```

```
kidneytest$rc[kidneytest$rc==""] <- NA
```

kidneytest\$rc[is.na(kidneytest\$rc)] # There are 16 blank spaces in kidney\$rc

#kidney\$rc is set to be chr, but will first changer to integer.

```
kidney$rc <- as.integer(kidney$rc)
```

#The mean will value will be used on these 16 blank spaces.

```
kidney$rc[kidney$rc==""] <- NA
```

#mean(kidney\$rc, na.rm = TRUE)) # This equals 4.385057, but will be rounded off

#to 4 to fit the other values in column "rc".

```
kidney$rc[is.na(kidney$rc)] <- 4
```

#Column "htn", "dm" and "cad"

```

#Kindneytest will be used to determinte how many blankspaces on all three.

#A visual look have them set to have them on the same rows.

#"htn"

kidneytest <- kidney

kidneytest$htn [kidneytest$htn==""] <- NA

kidneytest$htn [is.na(kidneytest$htn)] # 2 blank spaces.

#"dm"

kidneytest <- kidney

kidneytest$dm [kidneytest$dm==""] <- NA

kidneytest$dm [is.na(kidneytest$dm)] # 2 blank spaces.

#"cad"

kidneytest$cad [kidneytest$cad==""] <- NA

kidneytest$cad [is.na(kidneytest$cad)] # 2 blank spaces.

#Yes, they were on the same row and with only 2 blank spaces.

#Since the values are yes and no, the blank spaces will only be changed to NA
#and not given a mean value. The columns will have it class changed to factor
#and set a value of 1 to "yes", 2 to "no" and 3 to "NA".

#"htn"

kidney$htn <- as.factor(kidney$htn)

kidney$htn [kidney$htn==""] <- NA

kidney$htn <- factor(kidney$htn,levels=c("yes","no","NA"))

#"dm"

kidney$dm <- as.factor(kidney$dm)

kidney$dm [kidney$dm==""] <- NA

kidney$dm <- factor(kidney$dm,levels=c("yes","no","NA"))

#"cad"

kidney$cad <- as.factor(kidney$cad)

kidney$cad [kidney$cad==""] <- NA

kidney$cad <- factor(kidney$cad,levels=c("yes","no","NA"))

#In the next 3 columns, "appet", "pe" and ane

#It is the same situation as the previous 3 columns.

```

```

#This is shown by doing the kidneytest on the 3 columns

#“appet”
kidneytest <- kidney
kidneytest$appet [kidneytest$appet==""] <- NA
kidneytest$appet [is.na(kidneytest$appet)] # 1 blank spaces.

#“pe”
kidneytest <- kidney
kidneytest$pe [kidneytest$pe==""] <- NA
kidneytest$pe [is.na(kidneytest$pe)] # 2 blank spaces.

#“ane”
kidneytest$ane [kidneytest$ane==""] <- NA
kidneytest$ane [is.na(kidneytest$ane)] # 2 blank spaces.

#As in “htn”, “dm” and “cad”, these columns will be treated the same on the row
#with no values. Except for “appet” column that will have 1 as “good”, 2 as “poor”
#and 3 as “NA”.

#“appet”
kidney$appet <- as.factor(kidney$appet)
kidney$appet [kidney$appet==""] <- NA
kidney$appet <- factor(kidney$appet, levels=c("good", "poor", "NA"))

#“pe”
kidney$pe <- as.factor(kidney$pe)
kidney$pe [kidney$pe==""] <- NA
kidney$pe <- factor(kidney$pe, levels=c("yes", "no", "NA"))

#“ane”
kidney$ane <- as.factor(kidney$ane)
kidney$ane [kidney$ane==""] <- NA
kidney$ane <- factor(kidney$ane, levels=c("yes", "no", "NA"))

#The last column is classification”

#With sapply(kidney$classification, unique) it can be seen to have no blank spaces,
#and 2 values of “ckd” and “notckd”. This column will be changed from chr to a
#factor with 2 levels. 1 will be “ckd” and 2 will be “notckd”.

```

```
kidney$classification <- as.factor(kidney$classification)
```

```
kidney$classification <- factor(kidney$classification,levels=c("ckd","notckd"))
```

3

Explanation to every vector name in the data frame Kidney.

age = age

bp = blood_pressure

sg = specific_gravity

al = albumin

su = sugar

rbc = red_blood_cells

pc = pus_cell

pcc = pus_cell_clumps

ba = bacteria

bgr = blood_glucose_random

bu = blood_urea

sc = serum_creatinine

sod = sodium

pot = potassium

hemo = haemoglobin

pcv = packed_cell_volume

wc = white_blood_cell_count

rc = red_blood_cell_count

htn = Hypertension

dm = diabetes_mellitus

cad = coronary_artery_disease

appet = appetite

pe = peda_edema

ane = anemia

classification = Chronic Kidney Disease

This is a boxplot comparison graph, with that shows the the median number in red. The dots outside the boxes are extreme points.

```
library(ggplot2)
```

```
theme_set(theme_gray())
```

```
g <- ggplot(kidney, aes(age,classification))
```

```
g + geom_boxplot(varwidth=T, fill="darksalmon",colour="red") +
```

```
#warwidth=T for the boxes to be adjusted to the amount observations
```

```
#in the boxes.
```

```
labs(title="Comparison of age ",
```

```
  subtitle="Having cronic kidney disease or not ",
```


```
  caption="Kidney data frame",
```

```
  x="Age",
```

```
  y="White blood cell count")
```

```
  Comparison of age
```

```
  Having cronic kidney disease or not
```



Kidney data frame

This is an attempt to visualize either the difference or the similarity of a compound, 2 elements and a protein.

```
library(ggplot2)
```

```
ggplot(kidney, aes(y=age)) +
```

```
  geom_point(aes(x = al), color = "green")+
```

```
  geom_point(aes(x = su), color="blue")+
```

```
  geom_point(aes(x = pot), color="black")+
```

```
  geom_point(aes(x = rc),size=1, color="pink")+
```

```
  #adding 4 layers of vector to visualised with different color.
```

```
  coord_cartesian(ylim=c(10, 80), xlim=c(2.5, 6))+ #zooming in.
```

```
  labs(title="Overview = Albumin, Sugar, Potassium and Sodium",
```

```
        subtitle="To see any correlation to age",
```

```
        caption = "Kidney data frame",
```


```
        y= "Age",
```

```
        x= "Al(green) - Su(blue) - Pot(black) - Sod(pink)")
```

```
#last part is labeling the title, subtitle, caption, y and x.
```

```
Overview = Albumin, Sugar, Potassium and Sodium
```

```
To see any correlation to age
```



Kidney data frame

The amount of Hemoglobin in people that have normal red blood cells compared the abnormal.

```
library(ggplot2)
```

```
theme_set(theme_gray())
```

```
ggplot(kidney, aes(x=rbc, y=hemo)) +
```

```
  geom_bar(stat="unique", width=0.5, fill="olivedrab4") +
```

```
  labs(title="The sum of Hemoglobin",
```


```
        subtitle = "A comparison of how Hemoglobin is distributed among Red blood cells",
```

```
        caption="Kidney data frame",
```

```
        x="Red blood cells",y="Hemoglobin")
```

The sum of Hemoglobin

A comparison of how Hemoglobin is distributed among Red blood cells



Kidney data frame

This is a continue of the last graph and shows the amount of people having the similar value of hemoglobin.

```
library(ggplot2)
```

```
theme_set(theme_gray()) # adding grey background for clarity.
```

```
ggplot (kidney,aes(x=rbc,y=hemo))+
```

```
geom_violin(fill="deeppink4",width=1)+
```

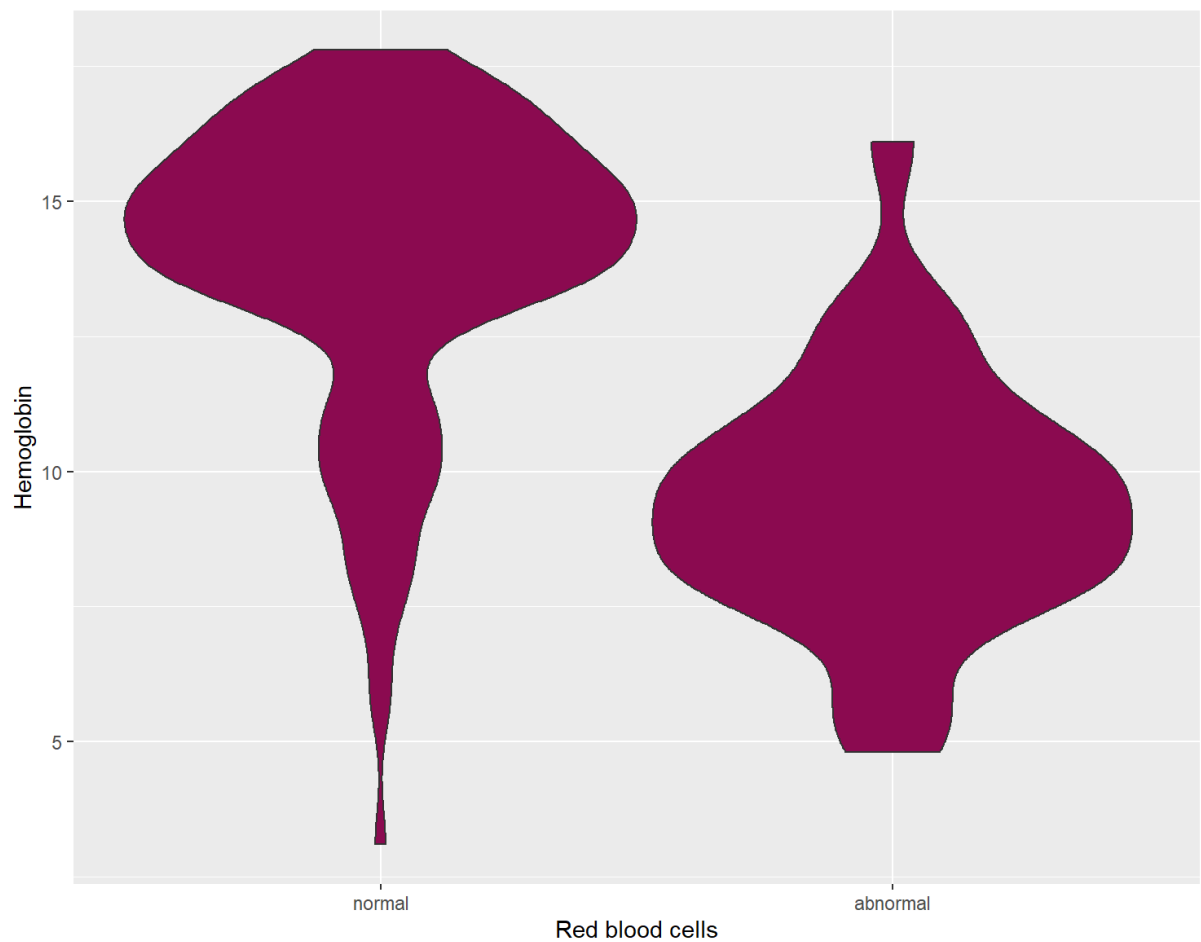
```
#width 1 for clearer visualisation.
```

```
labs(title="The count of Hemoglobin in each category",
```

```
caption="Kidney data frame",
```

```
x="Red blood cells",y="Hemoglobin")
```

```
The count of Hemoglobin in each category
```




Kidney data frame

bp = blood pressure. The graph shows blood pressure according to age.

```
qplot(age, bp, data = kidney, colour=bp, size=age)
```

#using the dataframe kidney, colour is set according to

#one of the parameter and size is set to the other.



This is a very zoomed in graph of potassium is getting lower the higher the value of hemoglobin is.

```
library(ggplot2)
```

```
g6 <- ggplot(kidney, aes(x=hemo, y=pot, colour=hemo)) +
```

```
  geom_point()+
```

```
  geom_smooth(method = "lm", fill="green")
```

```
#smooth emphasize the lower frequency behaviour.
```

```
#lm to make a line through the field to zoom in.
```

```
g6+coord_cartesian(ylim=c(2.4, 7.8), xlim=c(2.6,18))+
```

```
  labs(title="Correlation measure",
```


```
        subtitle="The average correlation between hemoglobin and potassium",
```

```
        caption="Kidney data frame",
```

```
        x="Hemoglobin",    y="Potassium")
```

```
  Correlation measure
```

```
  The average correlation between hemoglobin and potassium
```



Kidney data frame